

Language

English

Signed in as
admin

Log out

AI Job Market
2025-2026

Outputs-first
dashboard. New-data
processing calls the
same isolated offline
orchestrator used by
the release pipeline.

Data source

Evidence workspace

Baseline release

Upload &
> process a new
dataset

Workflow

- ☐ 1. Data Basic Clean
- ☐ 2. Data Ready for ML
- ☐ 3. AI Job Market Segmentation
- ☐ 4. Model Comparison
- ☒ 5. Best Model & Importance
- ☐ 6. Salary Prediction
- ☐ 7. Integrated Market Insight
- ☐ 8. Full Pipeline

Active evidence:

FPTCranes-PR32-main

5. Best model, diagnostics and uncertainty

Saved full model · frozen-configuration CV · historically scored test evidence

Report question

How does the saved full model generalize from temporal DEV validation to the historically scored test period?

Current takeaway: Test MAE decreased by \$156 versus final-configuration DEV CV.

Evidence: Final full configuration · 298 Test rows · ui-4ee3de59fabafb88

Main limitation: R² decreased by 0.025; the Test was already exposed to scoring and is not pristine.

Saved model

Random Forest · full 13

Test R²

0.803

Test MAE

\$14,877

Test MedAE

\$4,526

Test – CV R²

-0.025

Empirical q90

±\$42,254

298 historically scored rows · period 2026-03 · evidence ui-4ee3de59fabafb88

> How to read these metrics

Generalization Residual diagnostics Tuning Feature reliance **Trust and audit**

Question: How wide is the empirical historical-error reference, how often did it cover this same population, and what tail risk remains?

Empirical q90

±\$42,254

Same-population coverage

89.9%


Coverage denominator

298

RMSE / MedAE

6.6×

Test absolute-error distribution



Uncertainty conclusion

Finding: The ±\$42,254 historical-error band covers 89.9% of 298 scored rows.

Why it matters: RMSE is 6.6× MedAE, indicating that larger errors materially affect the average squared error.

Limit: Coverage was measured on the same historically scored population; it is not a calibrated future confidence interval.

Decision/use: Use the band as an empirical risk reference alongside the point estimate.

Three traceable historical benchmark examples

Record ID	Job title	Job category	Experience (years)	Actual salary (USD)	Full predi
historical_test:fad77ac0481dc34a38f33f3f28f28c326d5ef5c9f9187a69c87e5085d65e	Multimodal AI Engineer	Governance	8.0	\$145000	
historical_test:fad77ac0481dc34a38f33f3f28f28c326d5ef5c9f9187a69c87e5085d65e	AI Security Engineer	AI Engineering	4.0	\$234000	
historical_test:fad77ac0481dc34a38f33f3f28f28c326d5ef5c9f9187a69c87e5085d65e	AI Business Analyst	AI Engineering	3.0	\$264000	

The q90 band is computed and checked on the same Test population. It is not a formal confidence interval or guaranteed coverage for future or extrapolated scenarios.

Training logs

> Training log and audit downloads