

Language

English

Signed in as admin

Log out

AI Job Market 2025-2026

Outputs-first dashboard. New-data processing calls the same isolated offline orchestrator used by the release pipeline.

Data source

Evidence workspace

Baseline release

Upload & process a new dataset

Workflow

- 1. Data Basic Clean
- 2. Data Ready for ML
- 3. AI Job Market Segmentation
- 4. Model Comparison
- 5. Best Model & Importance
- 6. Salary Prediction
- 7. Integrated Market Insight
- 8. Full Pipeline

Active evidence:

FPTCranes-PR32-main

6. Controlled salary scenarios

Fixed Top-2 Random Forest · observed DEV policy · historical empirical uncertainty

Report question

How do validated job-category and experience scenarios compare under the fixed two-input salary model?

Current takeaway: The fixed two-input model records Test R² 0.821 with an empirical q90 absolute error of \$34,181.

Evidence: Inputs: job_category, years_of_experience · historical reference: 298 rows · ui-4ee3de59fabafb88

Main limitation: This is an academic scenario comparison, not a compensation recommendation or guaranteed salary range.

Model

Random Forest (fixed Top-2)

Active inputs

2

Test R²

0.821

Empirical q90

±\$34,181

Evidence ui-4ee3de59fabafb88 · fixed feature ablation, not independently tuned · no training occurs on this page

> How to read these metrics

1. Quick-load historical benchmarks

Multimodal AI Engineer · 8y · actual \$145,000

AI Security Engineer · 4y · actual \$234,000

AI Business Analyst · 3y · actual \$264,000

Examples were already included in Test scoring. They are traceable, not pristine or newly held out.

2. Controlled scenario builder

Job category

AI Engineering

Job title — observed DEV association

AI Agent Developer

Years of experience

5.00

Within observed DEV bounds: 2–12 years · title median 5 · pair support 29

Add validated scenario

3. Validation queue and actions

Scenario ID	Source	Job title	Job category	Experience (years)	Known actual (USD)
S001	benchmark	Multimodal AI Engineer	Governance	8.0	
S002	benchmark	AI Security Engineer	AI Engineering	4.0	
S003	benchmark	Multimodal AI Engineer	Governance	8.0	

Run batch salary prediction


Clear queue

4. Visual prediction results

Question: What salary range does this validated batch produce, how uncertain is the historical reference, and which rows have known actuals or extrapolation flags?

Scenarios 1–3 of 3

Salary predictions with empirical historical-error whiskers





Batch prediction conclusion

Finding: 3 scenarios range from 146,826 to 238,431 in predicted annual salary.

Why it matters: The batch has 3 known actuals; the empirical historical-error reference is $\pm 34,181$.

Limit: All scenarios use strict observed DEV bounds, but historical q90 still does not guarantee future coverage.

Decision/use: Compare controlled scenarios; do not treat point estimates or bands as compensation guarantees.

Whiskers use a same-population historical absolute-error q90. They are not formal confidence intervals or salary guarantees; lower bounds are clipped at zero.

☐ Extend growth curves to 0–15 years

Generate seniority growth view


Select scenarios / Job Title to display:

AI Security Engineer x

Multimodal AI Engineer x



Model scenario across years of experience



This model response curve is not a causal or necessarily monotonic salary trajectory.

Conclusion: Prediction of why AI salaries tend to decrease as years of experience increase => Senior engineers may not update their AI skills as quickly as younger generations (due to the rapid pace of technological innovation), leading to a decline in their competitive edge regarding the latest technologies.

Scenario ID	Job title	Job category	Experience (years)	Validation mode	Lower reference (USD)	Prediction (USD)	Upper reference (USD)	Known actual
S001	Multimodal AI Engineer	Governance	8.0	strict	\$112644	\$146826	\$181007	\$1
S002	AI Security Engineer	AI Engineering	4.0	strict	\$204250	\$238431	\$272612	\$2
S003	Multimodal AI Engineer	Governance	8.0	strict	\$112644	\$146826	\$181007	\$1

Download prediction audit v1

Download original enterprise-schema scenarios

The audit CSV uses signed variance. The original-schema file is an incomplete scenario interchange file; it does not place predictions into observed salary columns.